New dashboard Copy ▾

⟨ Zoom Out ⟩   ⊙ Last 1 hour   Refresh every 10s

irate of container_network_receive_bytes_total

irate container_network_transmit_bytes_total

rate of container_cpu_user_seconds_total

sum by (id, instance)(rate(container_cpu_usage_seconds_total[1m])) / count by (id, instance)(container_cpu_usage_seconds_total)

sum (rate (container_cpu_usage_seconds_total{container_label_com_docker_swarm_service_name="web_server"}[1m])) / sum (machine_cpu_cores) * 100

irate(container_cpu_usage_seconds_total{container_label_com_docker_swarm_service_name="web_server"}[5m])

irate of nginx_http_requests

irate(node_cpu{mode="idle"}[5m])

Grafana - New dashboard Copy

{cpu="cpu0",instance="10.0.0.32:9100",job="proxy_node-exporter",mode="idle"}

Number of web servers

0

nginx_metric_errors_total

nginx_metric_errors_total{instance="10.0.0.23:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.38:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.41:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.42:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.55:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.56:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.62:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.63:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.64:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.65:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.66:9145",job="web_server"} — nginx_metric_errors_total{instance="10.0.0.67:9145",job="web_server"}

+ ADD ROW